

DTV HW2- Detection

Introduction

In this assignment, it's objective is to modify a neural network to do detection. Dataset used is Pascal VOC 2007. The data has been split into 50% for training/validation and 50% for testing. The distributions of images and objects by class are approximately equal across the training/validation and test sets. In total there are 9,963 images, containing 24,640 annotated objects.

Experiment Setup

The neural network used in this assignment is faster-RCNN, using ResNet as its CNN layer. ResNet in this case has 101 layer implementation, edited from torchvision source code and adding additional module to try and improve its performance.

```
class Bottleneck(nn.Module):
    expansion = 4

    def __init__(self, inplanes, planes, stride=1, downsample=None):
        super(Bottleneck, self).__init__()
        self.conv1 = nn.Conv2d(inplanes, planes, kernel_size=1, stride=stride, bias=False) # change
        self.bn1 = nn.BatchNorm2d(planes)
        self.conv2 = nn.Conv2d(planes, planes, kernel_size=3, stride=1, # change
                               padding=1, bias=False)
        self.bn2 = nn.BatchNorm2d(planes)
        self.conv3 = nn.Conv2d(planes, planes * 4, kernel_size=1, bias=False)
        self.bn3 = nn.BatchNorm2d(planes * 4)
        self.relu = nn.ReLU(inplace=True)

        self.globalAvgPool = nn.AdaptiveAvgPool2d(1)
        self.fc1 = nn.Linear(in_features=planes*4, out_features=round(planes/4))
        self.fc2 = nn.Linear(in_features=round(planes/4), out_features=planes*4)
        self.sigmoid = nn.Sigmoid()

        self.downsample = downsample
        self.stride = stride
```

```
def forward(self, x):
    residual = x

    out = self.conv1(x)
    out = self.bn1(out)
    out = self.relu(out)

    out = self.conv2(out)
    out = self.bn2(out)
    out = self.relu(out)

    out = self.conv3(out)
    out = self.bn3(out)

    if self.downsample is not None:
        residual = self.downsample(x)

    original_out = out
    out = self.globalAvgPool(out)
    out = out.view(out.size(0), -1)
    out = self.fc1(out)
    out = self.relu(out)
    out = self.fc2(out)
    out = self.sigmoid(out)
    out = out.view(out.size(0), out.size(1), 1, 1)
    out = out * original_out



    out += residual
    out = self.relu(out)
    return out
```




Result

Graph For Reference result (only for 1 epoch due to long training time)

In the original github source code, they use pretrained model to speed up the process, however, due to the change in ResNet layer that will be implemented later, this process is skipped to obtain fair comparison.



Mean AP = 0.0032








Graph For Modified result (only for 1 epoch due to long training time)

Mean AP = 0.0043







Discussions

- Modifying ResNet directly affect the dimension of data when batch of data is considered
- Performance of modified result is better, however, we sacrificed slower testing time with not huge impact on training time.
- Due to long time in training and testing time, there is insufficient time to run the data for this report
- Modified codes differ mainly in `faster-rcnn.pytorch/lib/model/faster_rcnn/resnet.py` compared to reference

Performance

Loss function converges to zero much faster and with better mAP after only 1 epoch.